

HADOOP -----

Sudo su

Sudo apt-get update

Sudo apt-get install openjdk-8-jdk

Sudo java -version

Sudo update-alternatives --config java

Sudo addgroup Hadoop

Sudo adduser --ingroup Hadoop hduser

Sudo adduser hduser sudo

Sudo apt-get install ssh

Which ssh

Which sshd

Sudo su hduser

cd

Ssh-keygen -t rsa -P ""

Cat \$HOME/.ssh/id_rsa/pub >> \$HOME/.ssh/authorized_keys

Ssh localhost

Sudo mkdir -p /usr/local/Hadoop

Wget <https://archive.apache.org/dist/hadoop/common/hadoop-2.6.5/hadoop2.6.5.tar.gz>

tar xvzf Hadoop-2.6.5.tar.gz

cd hadop-2.6.5

sudo mv */usr/local/hadop

sudo chown -R hduser:hadoop /usr/local/hadoop

vim ~/.bashrc

#HADOOP VARIABLES START

export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64

export HADOOP_INSTALL=/usr/local/hadoop

export PATH=\$PATH:\$HADOOP_INSTALL/bin

export PATH=\$PATH:\$HADOOP_INSTALL/sbin

```
export HADOOP_MAPRED_HOME=$HADOOP_INSTALL
export HADOOP_COMMON_HOME=$HADOOP_INSTALL
export HADOOP_HDFS_HOME=$HADOOP_INSTALL
export YARN_HOME=$HADOOP_INSTALL
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_INSTALL/lib/native
export HADOOP_OPTS="-Djava.library.path=$HADOOP_INSTALL/lib" export
HADOOP_HOME_WARN_SUPPRESS=1
export HADOOP_ROOT_LOGGER="WARN,DRFA"
#HADOOP VARIABLES END
```

Source ~/.bashrc

Vim /usr/local/Hadoop/hadoop-env.sh

Export JAVA_HOME = /usr/lib/jvm/java-8-openjdk-amd64

sudo mkdir -p /app/hadoop/tmp

sudo chown hduser:hadoop /app/hadoop/tmp

vim /usr/local/hadoop/etc/hadoop/core-site.xml

```
<property>
    <name>hadoop.tmp.dir</name>
    <value>/app/hadoop/tmp</value>
    <description>A base for other temporary directories.</description>
</property>
```

```
<property>
<name>fs.default.name</name>
<value>hdfs://localhost:54310</value>
```

<description>The name of the default file system. A URI whose scheme and authority determine the FileSystem implementation. The uri's scheme determines the config property (fs.SCHEME.impl) naming the FileSystem implementation class. The uri's authority is used to determine the host, port, etc. for a filesystem.</description>

</property>

```
cp /usr/local/hadoop/etc/hadoop/mapred-site.xml.template  
/usr/local/hadoop/etc/hadoop/mapred-site.xml
```

```
vim /usr/local/hadoop/etc/hadoop/mapred-site.xml
```

<property>

<name>mapred.job.tracker</name>

<value>localhost:54311</value>

<description>The host and port that the MapReduce job tracker runs at. If "local", then jobs are run in-process as a single map

and reduce task.

</description>

</property>

```
sudo mkdir -p /usr/local/hadoop_store/hdfs/namenode
```

```
sudo mkdir -p /usr/local/hadoop_store/hdfs/datanode
```

```
sudo chown -R hduser:hadoop /usr/local/hadoop_store
```

```
vim /usr/local/hadoop/etc/hadoop/hdfs-site.xml
```

<property>

<name>dfs.replication</name>

<value>1</value>

<description>Default block replication.

The actual number of replications can be specified when the file is created. The default is used if replication is not specified in create time. </description>

</property>

<property>

<name>dfs.namenode.name.dir</name>

<value>file:/usr/local/hadoop_store/hdfs/namenode</value> </property>

<property>

<name>dfs.datanode.data.dir</name> <value>file:/usr/local/hadoop_store/hdfs/datanode</value>

</property>

hadoop namenode -format

hdfs namenode -format

start-all.sh

jps

if not working

stop-all.sh

sudo rm -rf /app/hadoop/tmp

sudo mkdir -p /app/hadoop/tmp

sudo chown hduser:hadoop /app/hadoop/tmp

sudo chown 750 /app/hadoop/tmp

start-all.sh

```
hdfs dfs -mkdir -p input_dir
```

```
cd
```

```
vim test.txt
```

```
hdfs dfs -put /home/hduser/test.txt input_dir
```

```
hdfs dfs -ls input_dir
```

```
hadoop jar /usr/local/hadoop/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.6.5.jar  
wordcount input_dir output
```

```
hdfs dfs -ls output
```

```
hdfs dfs cat output/part-r-00000
```

DOCKER

sudo apt-get update

curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -

sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu \$(lsb_release -cs) stable"

sudo apt-get update

apt-cache policy docker-ce

sudo apt-get install -y docker-ce

sudo service docker start

sudo docker run hello-world

i) ----- running an existing image (Grafana)

docker run -d -p 3000:3000 grafana/grafana-enterprise

sudo docker images

ii) -----creating a image

Mkdir DockerFiles

Cd DockerFiles

Touch Dockerfile

Vim Dockerfile

FROM ubuntu

MAINTAINER Nishanth VM Nishanth.codes@gmail.com

RUN apt-get update

CMD ["echo","hiiiiii everyone out there"]

Docker build -t Nishanth:1.1 .

Docker images

<<<get id of our image

Docker run <id of image>

Output: CMD

Hosting website -----

In ssh :::::::::::::::::::::

Sudo apt-get update

Sudo su

Sudo apt install apache2

cd..

cd var/www/html

rm index.html

touch index.html

vim index.html

<<<< write html,css, js code here >>>>

Sudo service apache2 start

In browser:::::::::::::::::::

<http://publicip/filename.html>

Migration (AMI)

Create AMI using right click

Go to AMI

Click add permission

Add user account ID (our ID)

SAVE

Go to create instance

Go to my amis / shared with me

Create

Reference

Scp command

Scp -i "keyname" filename publicdns